

# TP4 - Matrix transposition

Oguz Kaya  
oguz.kaya@universite-paris-saclay.fr

You can consult the Intel's webpage for AVX functions as well as their latencies for each architecture:  
<http://software.intel.com/sites/landingpage/IntrinsicsGuide>

It would be sufficient to filter for AVX, AVX2 and FMA type instructions on the leftside of the page when searching for an instruction.

To compile the program `program.cpp` with OpenMP and AVX/AVX2 instruction sets and generate the executable `program`, type the following command in the terminal:

```
g++ -O2 -std=c++11 -fopenmp -mavx2 -mfma program.cpp -o program
```

Part 1

## Parallel matrix transposition

In this exercise, we will develop efficient parallel kernels for transposing a dense matrix using AVX and OpenMP parallelism. The transpose of a square  $N \times N$  matrix  $A$  is a matrix  $B$  such that

$$B(i, j) = A(j, i)$$

for all  $1 \leq i, j \leq N$ .

We assume that  $N$  is sufficiently large ( $N = 1024, 2048, 4096, \dots$ ), a multiple of 32, and the datatype is `float`. The matrices are to be stored in a 1D array of size  $N^2$  using a row-major storage (i.e.,  $A(i, j) = A[i * N + j]$ ) You should implement your work in the provided skeleton code `mat-trans.cpp`.

*Ex. 1*

- Implement a sequential and scalar matrix transposition that computes Measure the performance in *GB/s* transposed (an  $N \times N$  matrix of floats takes  $4N^2$  bytes).
- Now implement the function `transAVX8x8_ps(_mm256 &line[8])` that takes as input a  $8 \times 8$  matrix stored line by line so that each line is in an AVX variable. The function should transpose this matrix, and overwrite `line[8]` with the transposed matrix. You should do this in three steps as explained in the class, using
  - `_mm256_unpacklo_ps` and `_mm256_unpackhi_ps` in the first step
  - `_mm256_shuffle_ps` in the second step
  - `_mm256_permute2f128_ps` in the final step.
- Now that you have the vectorized code to transpose a  $8 \times 8$  matrix, implement a version that makes use of this to compute  $B(i, j) = A(j, i)$ . Measure the performance.
- This time, implement an **in-place** transposition that computes the transposition on the matrix  $A$  directly, without using another matrix, i.e.,  $A(j, i) = A(i, j)$ . To do this, you will need to load, transpose, and store two  $8 \times 8$  tiles simultaneously. Measure the performance.